

CS 365 FINAL PROJECT: TRANSIT REGINA DATA WRANGLING

JANICE COTCHER

DECEMBER 5, 2025

Data Source: City of Regina Open Data Portal

Dataset: Transit Stops and Routes (November 20, 2025)

DATASET SOURCE & LICENSE

Source: [City of Regina Open Data Portal](#)

- Bus Stop Locations (yqrStops.json)
- Transit Routes (yqrRoutes.json)
- General Transit Feed Specification(routes.txt, stops.txt, trips.txt, stop_times.txt)

License: Open Government License - Regina

- Allows educational and commercial use
- No Personally Identifiable Information - only public infrastructure data

Why this matters: Understanding transit accessibility and route coverage for urban planning



Q Search

Datasets

Groups

Suggest a Dataset

Showcases

About

Contact Us

Organization



City of Regina

There is no description for this organization

Social

X Twitter

Facebook

License

Open Gov. License

Transit Network

The Regina Transit network that includes routes, stops and GTFS data.

Data and Resources

[Transit Network Live Map](#)
Map view of the Transit Network Routes and Stops data set. The map has been...

Explore ▾

[Transit Network Stops Geospatial SHP](#)
The geospatial transit stop locations in the City of Regina.

Explore ▾

[Transit Network Routes Geospatial SHP](#)
The geospatial transit route locations in the City of Regina.

Explore ▾

[Transit Network Routes and Stops Geospatial KML](#)
The geospatial transit route and stop locations in the City of Regina.

Explore ▾

[Transit Network Geospatial GTFS](#)
The Regina Transit GTFS data.

Explore ▾

[Transit Network JSON Data Service](#)
The live data feed for the Transit Network Routes and Stops, served in JSON.

Explore ▾

[Transit Network REST Data Service](#)
The live data feed for the Transit Network Routes and Stops, served in ArcGIS...

Explore ▾

[Transit Network SOAP XML Data Service](#)
The live data feed for the Transit Network Routes and Stops, served in SOAP.

Explore ▾

Bus GIS GTFS Network Open Data Route Schedule Stop Transit
Screenshot

TOOLS

VS CODE EXTENSIONS:

- Python
- Python Environments
- Data Wrangler
- Jupyter

TOOLS

VS CODE EXTENSIONS:

- Python
- Python Environments
- Data Wrangler
- Jupyter

RATIONALE

- large dataset[1]
- flexibility, memory management
- macOS
- personal expertise

[1] CS 365 - Data Cleaning: Concepts & Algorithms, Lecture 6: September 15, 2025

IMPORTED PYTHON LIBRARIES

```
import numpy as np # scientific computing for large, multi-dimensional arrays
from matplotlib import pyplot as plt # creates static, animated, and interactive visualizations
import pandas as pd #creates 2D, size-mutable, heterogeneous tables called data frames
import json # read and write json files
import plotly.graph_objects as go # interactive graphics like maps
from pyproj import Transformer # cartography and coordinate transformations
```

RAW DATA SNAPSHOT: BUS STOPS

```
raw_data > {} yqrStops20251120.json > [] features > {} 0 > {} attributes
1  {
2      "displayFieldName": "STOP_NAME",
3      "fieldAliases": { ... },
4      "geometryType": "esriGeometryPoint",
5      "spatialReference": {
6          "wkid": 26913,
7          "latestWkid": 26913
8      },
9      "fields": [ ... ],
10     "features": [
11         {
12             "attributes": {
13                 "OBJECTID": 59930,
14                 "ONSTREET": "University Park Dr",
15                 "ATSTREET": "Quance St (NB)",
16                 "LON": "-104.54913",
17                 "LAT": "50.44416",
18                 "STOP_ID": "0742",
19                 "STOP_NAME": "University Park Dr @ Quance St (NB)",
20                 "GLOBALID": "{05092908-E821-4704-86E0-4AA2BA573409}"
21             },
22             "geometry": {
23                 "x": 532013.8037999998,
24                 "y": 5588113.458399999
25             }
26         },
27         {
28             "attributes": {
29                 "OBJECTID": 59931,
30                 "ONSTREET": "University Park Dr",
31                 "ATSTREET": "Vic Square (NB)",
32                 "LON": "-104.54915",
33                 "LAT": "50.44592",
34                 "STOP_ID": "0743",
35                 "STOP_NAME": "University Park Dr @ Vic Square (NB)",
36                 "GLOBALID": "{1CA33E59-8211-476E-86F5-EDEB4E8F6F76}"
37             },
38             "geometry": ...
39         }
40     ]
41 }
```

```
# Load JSON data
try:
    with open('raw_data/yqrStops.json', 'r') as f:
        stop_data = json.load(f)
except json.decoder.JSONDecodeError as e:
    print("Invalid JSON", e)

# Normalize nested JSON structure
df_stops = pd.json_normalize(stop_data['features'])
```

DATA PROFILING: QUALITY ASSESSMENT - STOPS

- **Data types** - Are coordinates stored correctly?
- **Missing values** - Which columns have gaps?
- **Duplicates** - Any duplicate stop IDs?
- **Outliers** - Any stops in unexpected locations?
- **Cardinalities** - How many unique stops?

DATA PROFILING: QUALITY ASSESSMENT - STOPS

- **Data types** - Are coordinates stored correctly?
- **Missing values** - Which columns have gaps?
- **Duplicates** - Any duplicate stop IDs?
- **Outliers** - Any stops in unexpected locations?
- **Cardinalities** - How many unique stops?

STOP DATA TYPES

DATA PROFILING: QUALITY ASSESSMENT - STOPS

- **Data types** - Are coordinates stored correctly?
- **Missing values** - Which columns have gaps?
- **Duplicates** - Any duplicate stop IDs?
- **Outliers** - Any stops in unexpected locations?
- **Cardinalities** - How many unique stops?

STOP DATA TYPES

```
attributes.OBJECTID          int64
attributes.ONSTREET           object
attributes.ATSTREET           object
attributes.LON                object
attributes.LAT                object
attributes.STOP_ID            object
attributes.STOP_NAME          object
attributes.GLOBALID           object
geometry.x                   float64
geometry.y                   float64
dtype: object
```

MISSING VALUES

MISSING VALUES

Number of Missing Values: 1

Missing Values

attributes.OBJECTID	0
attributes.ONSTREET	0
attributes.ATSTREET	1
attributes.LON	0
attributes.LAT	0
attributes.STOP_ID	0
attributes.STOP_NAME	0
attributes.GLOBALID	0
geometry.x	0
geometry.y	0

dtype: int64

DUPLICATES

DUPLICATES

Number of Duplicate Stops - checked by stop ID: 0

CLEANING: BUS STOP TEXT STANDARDIZATION & TYPE CONVERSIONS

PROBLEMS IDENTIFIED

- Inconsistent text formatting (mixed case, whitespace)
- Coordinates stored as strings instead of numeric latitude and longitude
- Missing values in street names

Sample Solutions:

```
# Text standardization
df_stops['attributes.ONSTREET'] = df_stops['attributes.ONSTREET'].str.strip().str.
upper()
df_stops['attributes.ATSTREET'] = df_stops['attributes.ATSTREET'].str.strip().str.
upper()
...
# Missing value imputation
df_stops = df_stops.fillna({'attributes.ATSTREET': "DOROTHY ST (SB)"})

# Data correction
df_stops['attributes.ONSTREET'] = df_stops['attributes.ONSTREET'].str.replace(
    "1060 DOROTHY ST (SB)", "DOROTHY ST", regex=False
)
```

VISUAL CHECK OF BUS STOPS

RAW DATA SNAPSHOT: BUS ROUTES

```
{ } yqrRoutes.json X  
Users > janicecotcher > Documents > GitHub > CS365_presentation > raw_data > { } yqrRoutes.json > ...  
1  "displayFieldName": "ROUTE_NAME",  
2  "fieldAliases": {  
3    "OBJECTID": "OBJECTID",  
4    "SHAPE.LEN": "SHAPE.LEN",  
5    "ROUTE_NAME": "ROUTE_NAME",  
6    "ROUTE_NUM": "ROUTE_NUM",  
7    "ROUTE_ID": "ROUTE_ID",  
8    "SHAPE_ID": "SHAPE_ID",  
9    "ROUTE_COLOR": "ROUTE_COLOR",  
10   "ROUTE_TEXT_COLOR": "ROUTE_TEXT_COLOR"  
11 },  
12   "geometryType": "esriGeometryPolyline",  
13   "spatialReference": { "wkid": 26913, "latestWkid": 26913 },  
14   "fields": [  
15     { "name": "OBJECTID", "type": "esriFieldTypeOID", "alias": "OBJECTID" },  
16     {  
17       "name": "SHAPE.LEN",  
18       "type": "esriFieldTypeDouble",  
19       "alias": "SHAPE.LEN"  
20     },  
21     {  
22       "name": "ROUTE_NAME",  
23       "type": "esriFieldTypeString",  
24       "alias": "ROUTE_NAME",  
25       "length": 200  
26     },  
27     {  
28       "name": "ROUTE_NUM",  
29       "type": "esriFieldTypeString",  
30       "alias": "ROUTE_NUM",  
31       "length": 200  
32     },  
33     {  
34       "name": "ROUTE_ID",  
35       "type": "esriFieldTypeString",  
36       "alias": "ROUTE_ID",  
37       "length": 200  
38     }  
  ]  
Screenshot
```

```
{ } yqrRoutes.json X
Users > janicecotcher > Documents > GitHub > CS365_presentation > raw_data > {} yqrRoutes.json > ...
59 "features": [
60   {
61     "attributes": {
62       "OBJECTID": 9601,
63       "SHAPE_LEN": 184989.734238412,
64       "ROUTE_NAME": "RCMP - NORMANVIEW",
65       "ROUTE_NUM": "10",
66       "ROUTE_ID": "10-44",
67       "SHAPE_ID": "100009",
68       "ROUTE_COLOR": "FF0FF9",
69       "ROUTE_TEXT_COLOR": null
70     },
71     "geometry": {
72       "paths": [
73         [
74           [521564.20469999965, 5591834.3639000002],
75           [521562.4895999997, 5591802.6662000008],
76           [521562.23489999957, 5591795.5486999992],
77           [521562.19259999972, 5591788.5431999993],
78           [521562.43460000027, 5591781.4276999999],
79           [521562.88900000043, 5591774.4243000001],
80           [521563.48570000008, 5591767.3102000002],
81           [521564.3657999998, 5591760.3084999993],
82           [521565.6007000003, 5591753.3081999999],
83           [521566.90610000025, 5591746.4195000008],
84           [521568.4952999996, 5591739.5318999998],
85           [521570.29739999957, 5591732.6450999994],
86           [521570.36930000037, 5591732.4230000004],
87           [521579.8009999998, 5591699.1030999999],
88           [521582.8262000001, 5591688.1071000006],
89           [521594.2079999963, 5591646.4554999992],
90           [521596.58789999969, 5591637.1248000003],
91           [521598.61429999955, 5591627.4591000006],
92           [521600.42740000039, 5591617.9037999995],
93           [521601.8152000007, 5591608.2355000004],
94           [521602.4029000001, 5591603.3453000002],
95           [521603.5943, 5591589.6732000001],
96           [521603.98460000008, 5591581.0014999993]

```

Screenshot

OVER 90K LINES, ~4.9MB

```
{ } yqrRoutes.json ×
Users > janicecotcher > Documents > GitHub > CS365_presentation > raw_data > { } yqrRoutes.json > ...
1 [ ...
2   "displayFieldName": "ROUTE_NAME",
3   "fieldAliases": { ...
4     },
5   "geometryType": "esriGeometryPolyline",
6   "spatialReference": { "wkid": 26913, "latestWkid": 26913 },
7   "fields": [ ...
8     ],
9   "features": [ ...
10    ]
11 ]
12 ]
13 ]
14 ]
15 ]
16 ]
17 ]
18 ]
19 ]
20 ]
21 ]
22 ]
23 ]
24 ]
25 ]
26 ]
27 ]
28 ]
29 ]
30 ]
31 ]
32 ]
33 ]
34 ]
35 ]
36 ]
37 ]
38 ]
39 ]
40 ]
41 ]
42 ]
43 ]
44 ]
45 ]
46 ]
47 ]
48 ]
49 ]
50 ]
51 ]
52 ]
53 ]
54 ]
55 ]
56 ]
57 ]
58 ]
59 ]
60 ]
61 ]
62 ]
63 ]
64 ]
65 ]
66 ]
67 ]
68 ]
69 ]
70 ]
71 ]
72 ]
73 ]
74 ]
75 ]
76 ]
77 ]
78 ]
79 ]
80 ]
81 ]
82 ]
83 ]
84 ]
85 ]
86 ]
87 ]
88 ]
89 ]
90 ]
91 ]
92 ]
93 ]
94 ]
95 ]
96 ]
97 ]
98 ]
99 ]
90316 ]
90317 ]
90318 ]
```

LOADING ROUTES DATA

```
# Load routes data
with open('raw_data/yqrRoutes.json', 'r') as f:
    routes_data = json.load(f)

df_routes = pd.json_normalize(routes_data['features'])
```


LOADING ROUTES DATA

```
# Load routes data
with open('raw_data/yqrRoutes.json', 'r') as f:
    routes_data = json.load(f)

df_routes = pd.json_normalize(routes_data['features'])
```

	attributes.OBJECTID	attributes.SHAPELEN	attributes.ROUTE_NAME	attributes.ROUTE_N
0	9601	184989.734238	RCMP - NORMANVIEW	
1	9602	47401.585527	HARBOUR LANDING - UNIVERSITY	
2	9603	180172.679776	UNIVERSITY - SHERWOOD ESTATES	
3	9604	92096.531471	ALBERT S EXPRESS - ALBERT N EXPRESS	
4	9605	57978.961540	ARCOLA E EXP - ARCOLA DWTN EXP	
5	9606	115065.916376	GLENCAIRN - WHITMORE	
6	9607	664900.169679	DIEPPE/WESTERRA - BROAD NORTH	
7	9608	492246.798973	VARSITY PARK - MOUNT ROYAL	
8	9609	24622.479650	AIRPORT - DOWNTOWN	

9	9610	94968.614798	UNIVERSITY - ARCOLA EAST
10	9611	134424.603144	ARGYLE PARK - WOOD MEADOWS
11	9612	63697.887437	UNIVERSITY - ROCHDALE EXPRESS
12	9613	85020.879476	HILLSDALE - WALSH ACRES
13	9614	47980.063524	HAWKSTONE - LAKERIDGE
14	9615	108643.699393	MAPLE RIDGE WEST - MAPLE RIDGE EAST
15	9616	60200.279881	HERITAGE WEST - HERITAGE EAST
16	9617	90660.507444	PARKRIDGE - ALBERT PARK
17	9618	64226.528459	WESTHILL - ROSS INDUSTRIAL

DATA PROFILE: QUALITY ASSESSMENT - ROUTES

- **Data types** - Are coordinates stored correctly?
- **Missing values** - Which columns have gaps?
- **Duplicates** - Any duplicate route IDs?
- **Outliers** - Any routes in unexpected locations?
- **Cardinalities** - How many unique routes?

CLEANING: BUS ROUTE TEXT STANDARDIZATION & TYPE CONVERSIONS

DATA WRANGLER DEMO

BUS ROUTE CLEANING SUMMARY

- inconsistent text formatting
- missing colours
- hex values missing #
- standardize column names

BUS ROUTE CLEANING SUMMARY

- inconsistent text formatting
- missing colours
- hex values missing #
- standardize column names

Cleaned routes:22

	route_num	route_name	route_color
0	10	RCMP - NORMANVIEW	#FF0FF9
1	18	HARBOUR LANDING - UNIVERSITY	#80FF00
2	3	UNIVERSITY - SHERWOOD ESTATES	#A8A800
3	40	ALBERT S EXPRESS - ALBERT N EXPRESS	#00CECE
4	60	ARCOLA E EXP - ARCOLA DWTN EXP	#808000

PREPARE ROUTES FOR VISUALIZATION

```
# Create transformer to convert from UTM to lat/lon
transformer = Transformer.from_crs("EPSG:26913", "EPSG:4326", always_xy=True)
```

VISUAL CHECK OF A BUS ROUTE

TRANSFORMATION STEP 1: LOADING GTFS SCHEDULE DATA

GTFS (General Transit Feed Specification) provides detailed schedule information:

```
# Load GTFS files
stops_gtfs = pd.read_csv('raw_data/gtfs_data/stops.txt')
routes_gtfs = pd.read_csv('raw_data/gtfs_data/routes.txt')
trips_gtfs = pd.read_csv('raw_data/gtfs_data/trips.txt')
times_gtfs = pd.read_csv('raw_data/gtfs_data/stop_times.txt')

# Clean and standardize - sample
stops_gtfs['stop_name'] = stops_gtfs['stop_name'].str.upper().str.strip()
routes_gtfs['route_long_name'] = routes_gtfs['route_long_name'].str.upper().str.strip()
```

GTFS Data Loaded:

- 1400 stops
- 44 routes
- 5078 trips
- 285722 stop times

TRANSFORMATION STEP 2: PARSING DATE/TIME DATA

Problem: Time data stored as strings (HH:MM:SS)

Solution: Parse to datetime and derive time-based features

```
# Parse time columns
times_gtfs_clean['departure_datetime'] = pd.to_datetime(
    times_gtfs_clean['departure_time'],
    format='%H:%M:%S',
    errors='coerce'
)

times_gtfs_clean['departure_datetime'] = pd.to_datetime(
    times_gtfs_clean['departure_time'],
    format='%H:%M:%S',
    errors='coerce'
)

# Derive hour and minute features
times_gtfs_clean['arrival_hour'] = times_gtfs_clean['arrival_datetime'].dt.hour
times_gtfs_clean['arrival_minute'] = times_gtfs_clean['arrival_datetime'].dt.minute
times_gtfs_clean['departure_hour'] = times_gtfs_clean['departure_datetime'].dt.hour
times_gtfs_clean['departure_minute'] = times_gtfs_clean['departure_datetime'].dt.minute
```

Parsed time data and derived hour/minute features

	arrival_time	arrival_datetime	arrival_hour	arrival_minute	departure_datetime	departure_time
0	06:10:00	1900-01-01 06:10:00	6	10	1900-01-01 06:10:00	06
1	06:11:00	1900-01-01 06:11:00	6	11	1900-01-01 06:11:00	06
2	06:12:00	1900-01-01 06:12:00	6	12	1900-01-01 06:12:00	06
3	06:13:00	1900-01-01 06:13:00	6	13	1900-01-01 06:13:00	06
4	06:14:00	1900-01-01 06:14:00	6	14	1900-01-01 06:14:00	06
5	06:15:00	1900-01-01 06:15:00	6	15	1900-01-01 06:15:00	06
6	06:15:00	1900-01-01 06:15:00	6	15	1900-01-01 06:15:00	06
7	06:16:00	1900-01-01 06:16:00	6	16	1900-01-01 06:16:00	06
8	06:17:00	1900-01-01 06:17:00	6	17	1900-01-01 06:17:00	06
9	06:18:00	1900-01-01 06:18:00	6	18	1900-01-01 06:18:00	06

TRANSFORMATION STEP 3: MERGE/JOIN OPERATIONS

- 1400 stops
- Open Regina ASP.NET limit of 1000
- GTFS data contains some fields
- Imputation for the remaining

```
# Find stops in GTFS but not in geographic JSON data
missing_stops = stops_gtfs_clean[
    ~stops_gtfs_clean['stop_id'].isin(clean_stops['stop_id'])
]

# Merge datasets
merged_stops = pd.concat([clean_stops, missing_stops], ignore_index=True, sort=False)

# Impute missing street names from stop_name
for index, stop in merged_stops.iterrows():
    if pd.isna(stop['on_street']):
        merged_stops.at[index, 'on_street'] = stop['stop_name'].split(' @')[0]
    if pd.isna(stop['at_street']):
        merged_stops.at[index, 'at_street'] = stop['stop_name'].split('@ ')[:-1]
```

1224 new stops added

Found stops in GTFS data: 1224

Total stops after merge: 2224 (1224 added)

	object_id	on_street	at_street	lon	lat	stop_id	stop_name	id
0	59930.0	UNIVERSITY PARK DR	QUANCE ST (NB)	-104.54913	50.44416	0742	UNIVERSITY PARK DR @ QUANCE ST (NB)	{05E84AA2B/
1	59931.0	UNIVERSITY PARK DR	VIC SQUARE (NB)	-104.54915	50.44592	0743	UNIVERSITY PARK DR @ VIC SQUARE (NB)	{1C82EDEB4I
2	59932.0	FLEET ST	NORTH SERVICE RD (NB)	-104.549126	50.448609	0744	FLEET ST @ NORTH SERVICE RD (NB)	{AF5833DD54
3	59933.0	FLEET ST	FINES DR (NB)	-104.549111	50.449614	0745	FLEET ST @ FINES DR (NB)	{63706FC74F
4	59934.0	CAMBRIDGE AVE	MILFORD CRES (WB)	-104.5513	50.45041	0746	CAMBRIDGE AVE @ MILFORD CRES (WB)	{E6F4A928E

TRANSFORMATION: FEATURE DERIVATION

- GEOGRAPHIC REGIONS

Divide the city into four quadrants based on approximate city centre coordinates (Albert St & Victoria Ave Intersection)

TRANSFORMATION: FEATURE DERIVATION

- GEOGRAPHIC REGIONS

Divide the city into four quadrants based on approximate city centre coordinates (Albert St & Victoria Ave Intersection)

```
city_centre_lon = -104.618
city_centre_lat = 50.447
merged_stops_clean['region'] = ''

# Assign quadrants (NE, NW, SE, SW)
for stop in range(len(merged_stops_clean)):
    if float(merged_stops_clean['lat'].iloc[stop]) > city_centre_lat:
        if float(merged_stops_clean['lon'].iloc[stop]) > city_centre_lon:
            merged_stops_clean.at[stop, 'region'] = "NE"
        else:
            merged_stops_clean.at[stop, 'region'] = "NW"
    else:
        if float(merged_stops_clean['lon'].iloc[stop]) > city_centre_lon:
            merged_stops_clean.at[stop, 'region'] = "SE"
        else:
            merged_stops_clean.at[stop, 'region'] = "SW"
```

Derived regional classifications for all stops

Stop distribution by region:

- NW: 333 stops
- NE: 224 stops
- SW: 1403 stops
- SE: 264 stops

TRANSFORMATION 5: FEATURE DERIVATION - DISTANCE CALCULATIONS

- In ArcGIS/GIS systems, shape.len (shape length) represents the total length of the geometry in metres

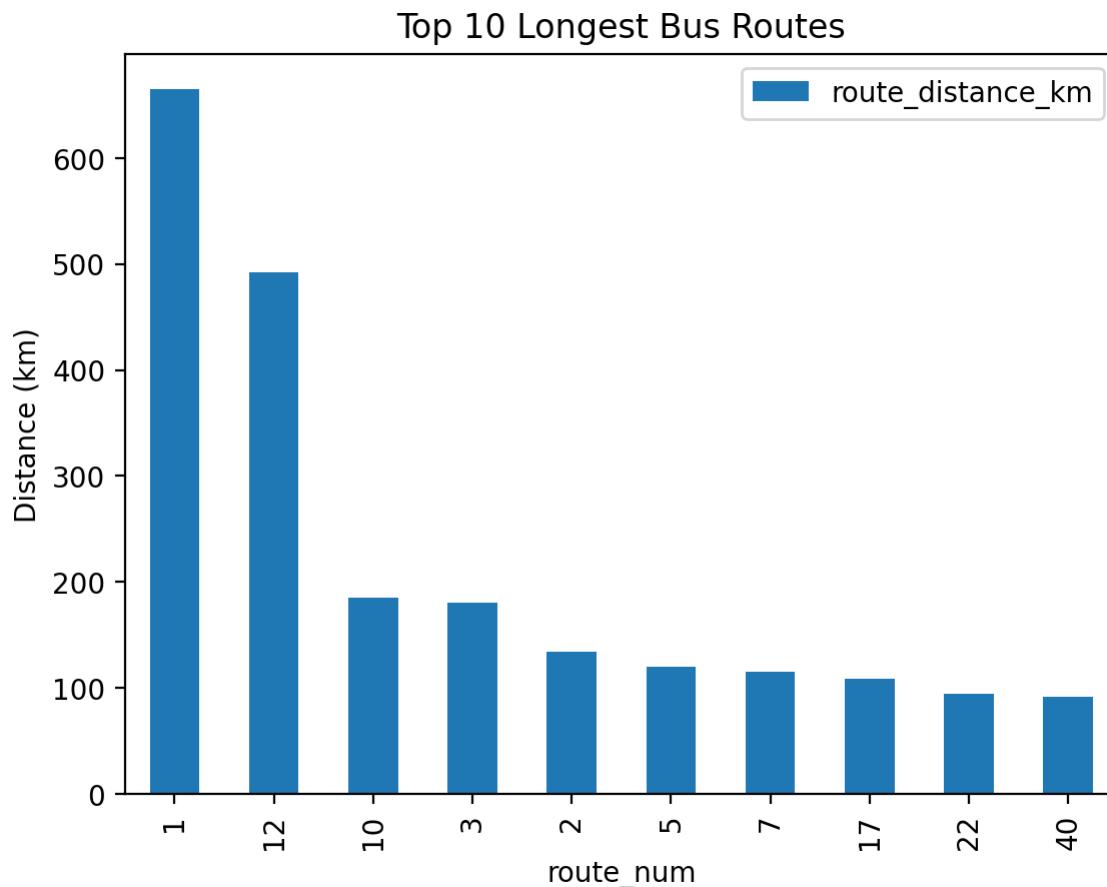
```
# convert shape_length into km
clean_routes['route_distance_km'] = clean_routes['shape_length'] / 1000
```

route_num	route_distance_km
0	10 184.989734
1	18 47.401586
2	3 180.172680
3	40 92.096531
4	60 57.978962
5	7 115.065916
6	1 664.900170
7	12 492.246799
8	24 24.622480
9	22 94.968615
10	2 134.424603
11	30 63.697887
12	4 85.020879
13	16 47.980064
14	17 108.643699
15	15 60.200280
16	9 90.660507
17	6 64.226528
18	8 89.851981
19	5 119.587762
20	21 68.741318

21

50

40.124363

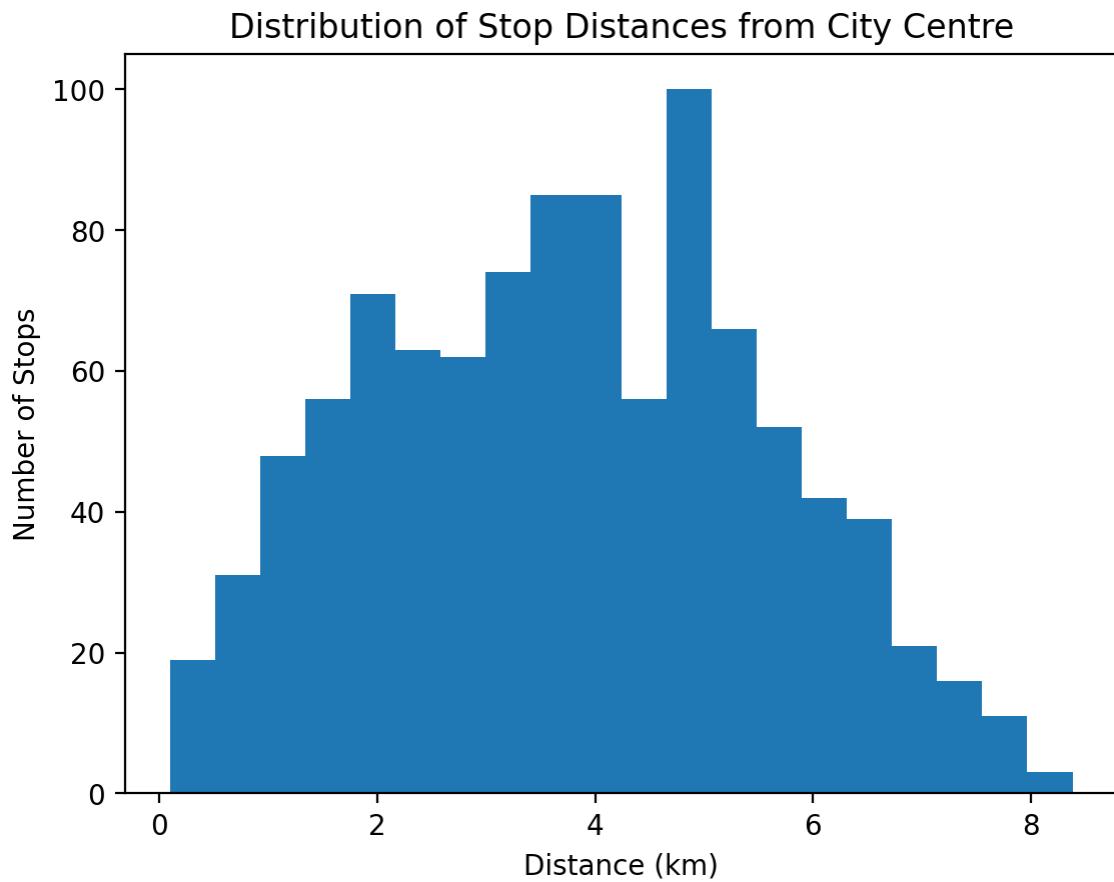


Calculate distance from city centre using coordinate geometry:

```
# Approximate conversion: ~111 km per degree latitude, ~85 km per degree longitude  
# at this latitude  
# Euclidean distance  
merged_stops_clean['distance_from_centre_km'] = np.sqrt(  
    ((merged_stops_clean['lat'].astype(float) - city_centre_lat) * 111)**2 +  
    ((merged_stops_clean['lon'].astype(float) - city_centre_lon) * 85)**2  
)
```

```
Distance statistics (km) :  
count    1000.000000  
mean      3.775980  
std       1.791905  
min       0.098964  
25%      2.319224  
50%      3.785285  
75%      5.068146  
max      8.381130  
Name: distance_from_centre_km, dtype: float64
```

Calculated distance from city centre for all stops



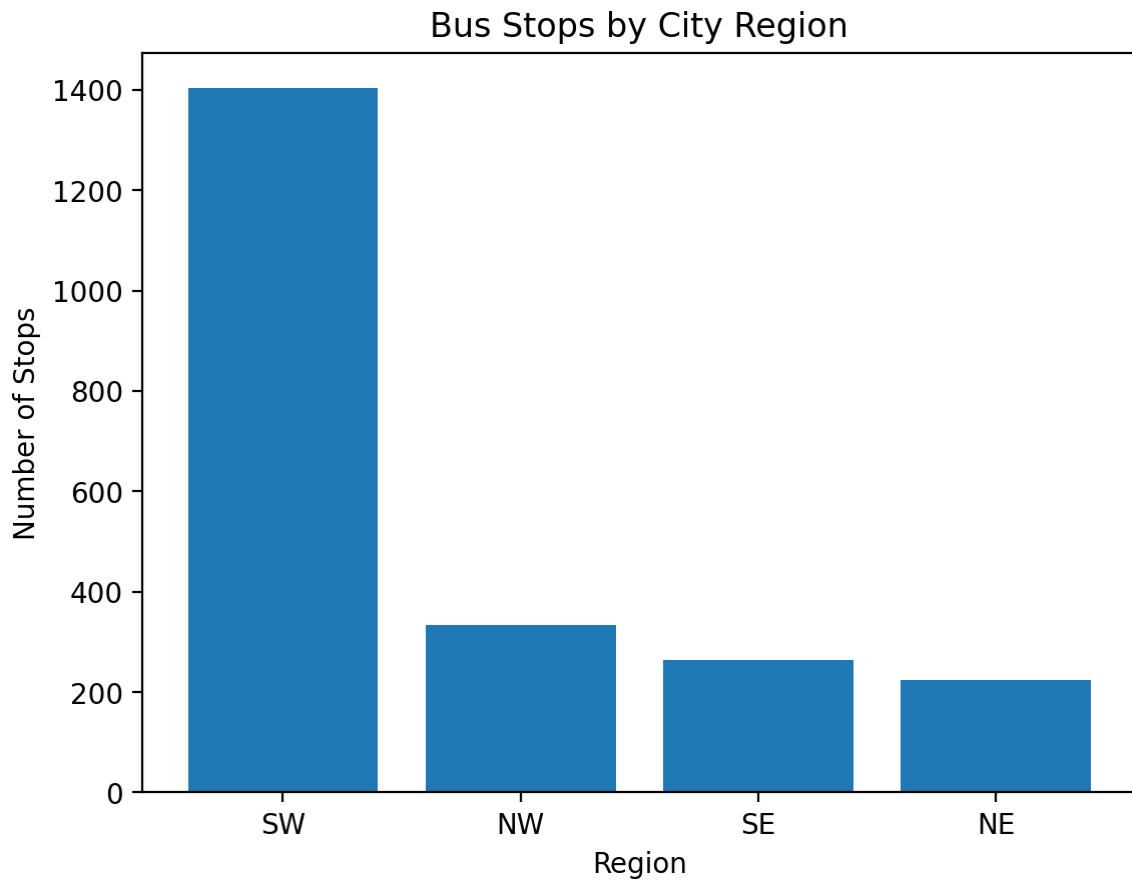
TRANSFORMATION 6: AGGREGATION

Aggregate stops by region to understand service distribution:

```
region_summary = clean_stops.groupby('region').agg({
    'stop_id': 'count',
    'distance_from_centre_km': ['mean', 'max']
}) .reset_index()
```

	stop_id	distance_from_centre_km	count	mean	min	max
region						
NE	224		4.113674	0.132092	7.650399	
NW	333		3.994900	0.306790	8.381130	
SE	264		3.926796	0.286489	7.536191	
SW	1403		2.723696	0.098964	5.336768	

Regional Summary Statistics:



TRANSFORMATION 7: RESHAPE (PIVOT)

Pivot analysis: How many stops does each route serve in each region?

```
# Join stop times → trips → stops to get route-region relationships
route_stops = (
    times_gtfs_clean
    .merge(trips_gtfs_clean[['trip_id', 'route_id']], on='trip_id')
    .merge(merged_stops_clean[['stop_id', 'region']], left_on='stop_id', right_on=
'stop_id')
    .groupby(['route_id', 'region'])
    .size()
    .reset_index(name='stop_count')
)

# Pivot to wide format
route_region_pivot = route_stops.pivot(
    index='route_id',
    columns='region',
    values='stop_count'
).fillna(0)
```

Routes by Region (Pivoted):

region	NE	NW	SE	SW
route_id				
1-44	0.0	0.0	0.0	11105.0
1-45	0.0	0.0	0.0	10997.0
10-44	222.0	1966.0	0.0	6468.0
10-45	210.0	1862.0	0.0	6124.0
12-44	0.0	1400.0	1536.0	8978.0
12-45	0.0	1400.0	1536.0	8978.0
15-44	28.0	0.0	316.0	766.0
15-45	28.0	0.0	316.0	766.0
16-44	0.0	56.0	0.0	388.0
16-45	0.0	56.0	0.0	388.0

PROFILING: OUTLIERS & CARDINALITIES

Found 0 outlier stops

Outlier Detection: Identify stops unusually far from city centre

Stops beyond 15km from city centre:

Unique stops: 2224

Unique routes: 44

Unique regions: 4

Unique trips: 5078

Cardinality Analysis: Count unique values in key dimensions

BEFORE/AFTER EVIDENCE

QUANTITATIVE COMPARISON OF DATA QUALITY IMPROVEMENTS:

Stop rows: 1000

Route rows: 22

Missing ATSTREET: 1

Coordinate type: object (string)

Stop ID type: object (string)

Features: 8 columns

Stop rows: 2224 (+1224 from GTFS)

Route rows: 22

Missing ATSTREET: 0

Coordinate type: object (numeric-ready)

Stop ID type: object (int32)

Features: 20 columns

New derived features: region, route_distance_km, distance_from_centre_km

Parsed time features: arrival_hour, arrival_minute, departure_hour, departure_minute

BEFORE (RAW DATA)

AFTER (CLEANED & TRANSFORMED)

SUMMARY OF TRANSFORMATIONS

COMPLETED TRANSFORMATIONS (8 OPERATIONS ACROSS 5 CATEGORIES):

1. Type Fixes & Parsing

- Converted stop_id from string to int32
- Parsed arrival/departure times to datetime
- Derived hour and minute features

2. Text Cleanup

- Stripped whitespace from all text columns
- Converted to uppercase for consistency
- Fixed malformed addresses

3. Missing Data Handling

- Imputed missing ATSTREET values
- Generated street names from stop names for GTFS stops

4. Join/Merge

- Merged geographic stops with GTFS schedule data
- Joined stop times → trips → routes → stops

5. Feature Derivation

REPRODUCIBILITY

HOW TO REPRODUCE THIS ANALYSIS:

1. Install dependencies:

```
pip install pandas plotly pyproj numpy jupyter
```

2. Directory structure:

```
project/
    ├── presentation.ipynb
    ├── transit_data.ipynb
    └── raw_data/
        ├── yqrStops.json
        ├── yqrRoutes.json
        └── gtfs_data/
            ├── stops.txt
            ├── routes.txt
            ├── trips.txt
            └── stop_times.txt
    └── README.md
```

3. Run notebook:

- Command line: `jupyter notebook transit_data.ipynb`
- In VS Code: Execute all cells sequentially
- Or run with Mercury: `mercury run presentation.ipynb`